

## Matrix Multiplication ijk Forms with Open MPI

I used the same method of partitioning for each of the three forms. My initial method that I started to implement was to broadcast the A matrix to each processor and then scatter portions of the B matrix to each process, and then from there I can do the partial matrix multiplication on each process and send the results back to process 0 with a gather command. When I started implementing this method, I realized that it would be very complicated to scatter columns of B to each processor (unless I read it in by column major). So, I changed my method of partitioning and decided I would broadcast all of B to each process and then scatter rows of A to each process. This was a simpler method because if the programmer follows the scatter command documentation, it is easy to send contiguous array values (rows) without any extra work that the programmer needs to do (such as reading in a matrix in column major order). Here are the steps for how my program works:

- 1.) Create a derived data type for the "Details" object that holds the matrix form (ijk, kij, ikj), the flag (I, R) and the n value.
- 2.) If my\_rank is 0, then read in the matrix form, the flag and the n value into the details object...then read in the two matrices in row major order. Next, start the timer.
- 3.) Broadcast the initialized details object from processor 0 to all the processors in the comm world.
- 4.) Now that every processor has the value of n (through the details object), they use this value to allocate the appropriate amount of space for their matrices.
- 5.) Scatter  $(N*N) / \text{comm\_sz}$  matrix A values from processor 0 to all the other processors.
- 6.) Broadcast matrix B values from processor 0 to all the other processors.
- 7.) Calculate local matrix multiplications in each processor.
- 8.) Gather the final matrix by sending to processor 0.
- 9.) If my\_rank == 0, stop the timer.

This method ended up being successful in reducing the timing of the work as the number of processors increased. After I did my runs using this approach, I realized something interesting. I found out that the "IJK" form was the slowest. When I began to look at the for loop for this form, I realized that at the innermost loop, we are accessing values in matrix B by column. This is important because this means that each index into the matrix is 4800 slots away from the previous index. This causes many cache misses and wasted clock cycles. I was able to notice this fairly easily because I chose to index my matrices using multiplication instead of brackets (which helped me easily visualize what the CPU was really doing). After I came to this realization, I soon came to an epiphany...that if I was to read in Matrix B via column major order and change my innermost for loop for IJK to access B by using this line of code:

" $(*(B + (j * n) + k))$ ", I would get the best possible cache performance out of any of the other forms because I would be looping through by row in all 3 matrices (the resulting matrix that we are writing to ( C ), the first matrix ( A ), and the second matrix ( B ) ).

```
(* (C + (i * n) + j)) += (* (A + (i * n) + k)) * (* (B + (j * n) + k));
```

The fastest method with my current implementation was the IKJ form. This form as well as the KIJ form are very similar in cache efficiency for the right hand side of their innermost loop, but the left hand side for IKJ is more efficient because it only indexes a new column in the resulting matrix ( C ) N times while the KIJ matrix does this  $N^2$  times. This operation is inefficient for the caches and that's why IKJ was a tad bit faster.


I also began to think about how I would handle the situation if the work was NOT divided up evenly like it was in this assignment. I did some research and found that ScatterV takes care of this issue for us, so I would have used that along with GatherV.

In the end, we can conclude that the method you use for matrix multiplication does actually matter. Initially, I did not think performance differences would be so major between the different IJK forms, but this assignment taught me otherwise. It is easy to see that there is a major difference if we simply look at the timing runs from 1 processor running IJK (min = 1249.0245 seconds) vs IKJ (min = 928.05212) vs KIJ (min = 931.24939). The difference in timing is about 5 minutes, which is huge.


For Loop Form	Number of Processors (Times are in Seconds)					
	1	4	8	12	16	20
IJK	-> 1249.0245 <- 1250.2384 1249.1558 1253.6686 1267.586	323.3315 322.20503 -> 315.8317 <- 316.21107 318.46949	161.19017 158.82207 160.07774 158.5909 -> 158.56746 <-	-> 107.17693 <- 107.22562 107.25106 107.42812 107.80458	81.387831 -> 81.301377 <- 82.161611 81.610694 82.23929	66.405539 66.367359 66.218579 -> 65.577977 <- 65.7224
IKJ	928.55334 928.16964 928.27876 -> 928.05212 <- 928.15172	234.65949 -> 234.47495 <- 234.69507 234.60153 234.6494	118.81671 -> 118.77531 <- 118.80128 118.89058 118.94237	80.37642 -> 80.118363 <- 80.299319 80.150053 80.447911	61.008533 60.965482 61.029474 61.079599 -> 60.914421 <-	49.400135 49.619495 -> 49.396811 <- 49.397762 49.409853
KIJ	931.31674 932.4446 -> 931.24939 <- 931.66506 931.53473	234.85126 235.55411 235.47941 -> 234.8344 <- 235.18775	118.92132 119.24383 -> 118.90616 <- 119.26628 119.19419	80.606154 80.61748 80.644684 -> 80.596913 <- 80.617037	-> 61.127414 <- 61.25233 61.300891 61.263628 61.247006	49.660248 49.622691 49.539667 49.499813 -> 49.444644 <-

For Loop Form	Number of Processors (E = Efficiency, SU = SpeedUp)					
	1	4	8	12	16	20
IJK	E: 1.0 SU: 1.0	E: 0.9886788596584826 SU: 3.9547154386339303	E: 0.9846160271470578 SU: 7.876928217176462	E: 0.9711546598694327 SU: 11.653855918433193	E: 0.9601809235039155 SU: 15.362894776062648	E: 0.9523200906304261 SU: 19.046401812608522
IKJ	E: 1.0 SU: 1.0	E: 0.9895002856381886 SU: 3.9580011425527544	E: 0.9766888000544893 SU: 7.813510400435915	E: 0.9652927714794505 SU: 11.583513257753406	E: 0.9522089605021444 SU: 15.23534336803431	E: 0.9393846497499606 SU: 18.78769299499921
KIJ	Efficiency: 1.0 SU: 1.0	E: 0.991389453589423 SU: 3.965557814357692	E: 0.9789751325751331 SU: 7.831801060601065	E: 0.9628670993060656 SU: 11.554405191672787	E: 0.9521601367759479 SU: 15.234562188415167	E: 0.9417090655966701 SU: 18.834181311933403


IKJ SpeedUp




IKJ Efficiency




KIJ SpeedUp



KIJ Efficiency



LJK SpeedUp



LJK Efficiency

